

# Collaborative Randomized Incremental Constructions on the Example of N-Queens

## Presentation of Bachelor Results

by Tristan Menke, Supervised by Florian Fey

Westfälische Wilhelms-Universität Münster

May 10, 2020

# Introduction

- ▶ Many computational tasks can be stated as Constraint Satisfaction Problems(CSPs)
  - ▶ Many AI tasks
  - ▶ Scheduling of processes and traffic
  - ▶ Many games like Sudoku
- ▶ I present an easy to implement, fast way to solve many problems, focused mainly on CSPs
- ▶ Finding fast, general and easy to implement solutions for these problems is useful in many areas
- ▶ General approaches may not be optimal for any one problem
- ▶ ...but if they are fast, that's good enough in practice

# Constraint Satisfaction Problems(CSPs)

A CSP is:

- ▶ A set of variables  $X_1, \dots, X_n$  with domains  $D_1, \dots, D_n$
- ▶ A set of constraints  $R_{i,j} \subseteq D_i \times D_j$  (or with more variables), that restrict which values the variables may take.
  - ▶ Example:  $R(X_1, X_2, X_3)$  specifies that  $X_1 + X_2 = X_3$  must be true
- ▶ A solution to a CSP is an assignment of all variables such that all constraints for each variable are satisfied
- ▶ Generally CSPs are in NP-Complete, but specific problems may be easier to solve
- ▶ I use the N-Queens problem as a specific CSP to implement all approaches presented

# The N-Queens Problem as a CSP

The N-Queens problem consists of:

- ▶ An  $N \times N$  chess board with the goal to place  $N$  queens
- ▶  $N$  variables that specify the queen positions, with values in  $\{1, \dots, N\}$
- ▶ 3 constraints on each variable:
  - ▶ No queen may be on the same **row** as another queen
  - ▶ No queen may be on the same **column** as another queen
  - ▶ No queen may be on the same **diagonal** as another queen
- ▶ We need to eliminate as many constraints as possible to quickly check the validity of an assignment

# Eliminating constraints

- ▶ Eliminate the row constraint by interpreting the indices of variables as row positions, their values as column positions
  - ▶ All queen positions can be fully specified by an array **queens** of integers with length  $N$
  - ▶  $\text{queens}[5] = 6$  means that the queen on row 5 is in column 6
- ▶ Eliminate the column constraint by making sure that the variable values are a permutation of  $(1, \dots, N)$  at all times
  - ▶ Initialize the array as  $(1, \dots, N)$  and only ever change it by swapping the values of two elements within the array
- ▶ Only the diagonal constraint remains to be checked

## An observation about diagonals

- ▶ The sum of row and column indices is constant on positively sloped diagonals, the difference is constant on negatively sloped diagonals

Sum:

0	1	2
1	2	3
2	3	4

Difference:

0	-1	-2
1	0	-1
2	1	0

# Checking the Diagonal Constraint

- ▶ A simple observation about diagonals:
  - ▶ There are  $2*N-1$  diagonals with positive and negative slope each
  - ▶ All positions on the same diagonal with positive slope have the same value for  $row + col$
  - ▶ ... and  $row - col$  on each diagonal with negative slope
- ▶ The diagonals that a queen threatens can be calculated as  $row + queens[row]$  and  $row - queens[row] + N - 1$  (adding  $N-1$  to always have positive indices)
- ▶ Use two arrays of Boolean values with length  $2*N-1$  to save which diagonals are occupied
- ▶ Checking new queens is done in constant time by calculating their diagonal indices

# Backtrack Search

- ▶ Simple and general, but inefficient algorithm used to solve CSPs
- ▶ Incrementally append new variable assignments of variable  $X_{i+1}$  to a sub-sequence  $(X_1, \dots, X_i)$  until all variables are set
- ▶ Check all constraints after each step, backtrack and change previous assignments if no further extensions can be made
- ▶ The simplest version iterates through all variables in a deterministic order, and tries each possible extension for that variable in a deterministic order
- ▶ Terminate if all variables are set or first variable runs out of choices



# Backtrack Search pseudo-code

Backtrack search is usually implemented recursively


```
1 bool backtrack(int row)
2 {
3     if (row == N)
4         return true;                // all queens placed
5
6     // try choices left-to-right
7     for each valid column c in row
8     {
9         place queen in column c;
10        if (backtrack(row+1))        // next row recursively
11            return true;            // subproblem solved!
12        else
13            remove queen in column c; // try next solution...
14    }
15
16    return false;                    // no solution found
17 }
```

Detailed implementation notes are in the thesis


## The 'search' part of Backtrack Search

- ▶ Backtrack Search can be interpreted as a depth-first search algorithm through a tree of all possible sequences of assignments
- ▶ The goal of search is to find a node with depth  $N$  (counting the root as depth 0)
- ▶ The leftmost node is searched first. If a node has no children, but is not depth  $N$ , the search algorithm "backtracks" to the previous node and checks the next child
- ▶ Backtrack Search is extremely slow for big problems, due to its exponential nature


## Backtrack Search deterministic order




## Backtrack Search deterministic order




## Backtrack Search deterministic order




## Backtrack Search deterministic order



## Backtrack Search deterministic order




## Backtrack Search deterministic order







## Backtrack Search deterministic order



## Backtrack Search deterministic order



## Backtrack Search deterministic order



# Randomized Incremental Constructions(RIC)

- ▶ Another general approach for solving CSPs
- ▶ Place a queen in each row in a random valid location
- ▶ Two central operations:
  - ▶ Choose a random possible extension for a partial solution through an **increment** operation
  - ▶ **Reset** to an empty sequence if no more extension can be made
- ▶ No attempt is made to improve a failed partial solution
- ▶ Simple to implement for specific CSPs by defining the operations
- ▶ Somewhat counter-intuitively, this is much faster in practice than Backtrack Search

# RIC operations increment and reset


- ▶ **Reset** operation

- ▶ ... restores the solution back to the empty state
- ▶ ... is trivial for N-Queens


- ▶ **Increment** operation

- ▶ ... computes all possible extensions for the solution
- ▶ ... chooses one possible extension at random
- ▶ ... returns true only if there was a valid extension that has been chosen


## RIC randomized search



## RIC randomized search




## RIC randomized search







## RIC randomized search



## RIC randomized search



## RIC randomized search



# RIC pseudo-code

The goal is to increment N times without failure.

```
1 void RIC(int N)
2 {
3     solution=empty solution
4     while(!success)                //Repeat until success
5     {
6         for row=1 to N              //increment N times
7         {
8             if(!increment(solution, row)) //If increment failed
9             {
10                 reset(solution);        //Reset and repeat
11                 break;
12             }
13         }
14         if(row == N)                  //All queens placed
15             success = true;
16     }
17 }
```

## Increment pseudo-code

```
1 bool increment(solution, int row)
2 {
3     possible_extensions = [];
4     for each column c in row //Compute possible extensions
5     {
6         calculate diagonal indices row+c and row-c;
7         if(both diagonals are unoccupied)
8             add c to possible_extensions;
9     }
10    if(possible_extensions is not empty)
11    {
12        c = random element of possible_extensions;
13        place queen in row at column c;
14        return true; //Solution has been extended
15    }
16    return false; //No extension can be chosen
17 }
```


## Advantages of RIC over Backtrack Search

- ▶ Cheap loop instead of recursion
- ▶ Does not need to keep track of previous choices
- ▶ Bad paths are quickly reset, not explored thoroughly like for Backtrack Search
- ▶ Worst-case rarely occurs

# A comparison of Backtrack Search and RIC

- ▶ Backtrack Search may take a long time to find a solution if it traverses the tree in a bad order
- ▶ The traversal-order has been randomized and sampled 10000 times for each data-point
- ▶ We compare the steps required for problems of various board sizes
  - ▶ A step is an increment for RIC and a recursive step for Backtrack Search
  - ▶ This allows us to compare small time-frames accurately
  - ▶ ... and prevent interference of outside factors
- ▶ Backtrack Search performs much worse and has worse scaling behavior

# A comparison of Backtrack Search and RIC



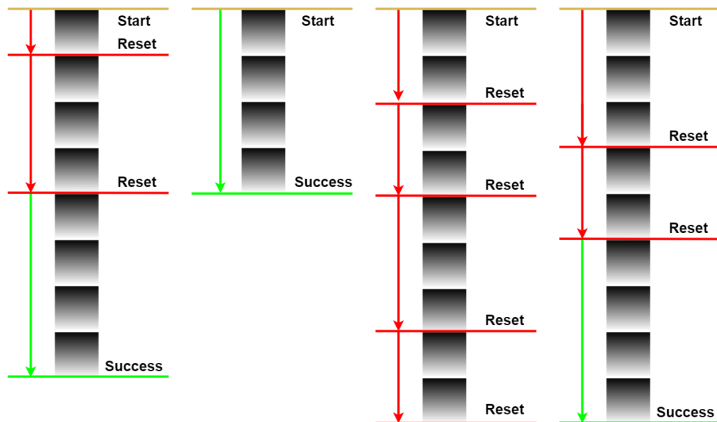


## Parallel Execution of RIC

- ▶ RIC can be run in parallel due to its randomized nature
- ▶ Different threads make distinct random choices
- ▶ Terminate after one thread finds a correct solution
- ▶ No parallelisation of the workload, speedup exclusively due to faster convergence
- ▶ Early termination through a single shared variable is extremely cheap
- ▶ Implementation for the CPU is very straight forward in a parallel block with a check for success of other threads at each step
- ▶ All threads operate independently, so no restrictions of freedom are needed

# Variance of steps across threads

Different threads making distinct random choices leads to variance in the required steps




# Collaborative RIC

- ▶ RIC implicitly assumes that any partial solution, regardless of how many steps have been taken, has the same probability to succeed in the end
- ▶ Instead, we assume that solutions with some progress have an inherently better chance to have at least one valid solution
- ▶ Starting a high number of threads has a high chance of at least one of them having a correct solution in their explored branch
- ▶ On this assumption, it makes no sense to reset any solution completely, if there are other solutions still active
- ▶ If there are still active threads, we replace the reset operation with a **copy** operation
- ▶ The copy operation chooses a random active thread and copies the solution of that thread to the calling thread

# Solution sharing in collaborative RIC

- ▶ Copied solutions are likely to choose different extensions
- ▶ Promising branches of the tree are explored more thoroughly
- ▶ ... but still reset after at most N steps
- ▶  $\Rightarrow$  We combine the thoroughness of exploring more paths from Backtrack Search with the quick resets of RIC



## Collaborative RIC pseudo-code

```
1 void collab_RIC(int N)
2 {
3     execute in parallel with shared variable success
4     {
5         solution=empty solution
6         while(!success)    //Repeat until success
7         {
8             for row=1 to N    //increment N times
9             {
10                if(success)    //Another thread succeeded
11                    break;
12
13                if(!increment(solution, row)
14                    && !copy(solution))    //We cannot copy
15                {
16                    reset(solution);    //Reset and repeat
17                    break;
18                }
19            }
20            if(row == N)    //All queens placed
21                success = true;
22        }
23    }
```

## Copy operation

- ▶ Copy operation iterates through all threads in random order
  - ▶ ... checks if that thread currently has an active solution
  - ▶ ... and copies that solution if it does
  - ▶ ... or returns false if there was no thread with an active solution
- ▶ The copy operation quickly finds an active thread when almost all threads are active
- ▶ ... which is almost always because the copy and reset operation are fast and we start with all threads being active

## Copy: restricting freedom

- ▶ The copy operation depends on the state of another thread's solution
- ▶ If that solution is changed while being copied, the copying thread receives a broken solution
- ▶ Two main possibilities to restrict the freedom of threads:
  - ▶ Barriers
    - ▶ ensure that all threads finish their increment operations before any thread attempts to copy
    - ▶ force all threads to synchronize
    - ▶ are very slow, especially for high thread counts
  - ▶ Locks
    - ▶ define mutually exclusive operations that can not occur concurrently
    - ▶ limit only the freedom of threads that are affected
    - ▶ are much faster and have better scaling for high thread counts
- ▶ Locks are more well suited for this task

## Copy: Locks

- ▶ A lock defines a critical region of code
- ▶ Each thread that wants to execute this critical region needs to acquire ownership of the lock
- ▶ Only one thread may own the lock at the same time
- ▶ All other threads need to wait for the lock to be released
- ▶ We assign a lock to each thread
  - ▶ The thread acquires its own lock when calling increment and reset
  - ▶ The thread acquires the other thread's lock when copying that thread's solution
  - ▶  $\Rightarrow$  A solution can never be changed while being copied from
  - ▶ Threads can run independently unless they are being copied from

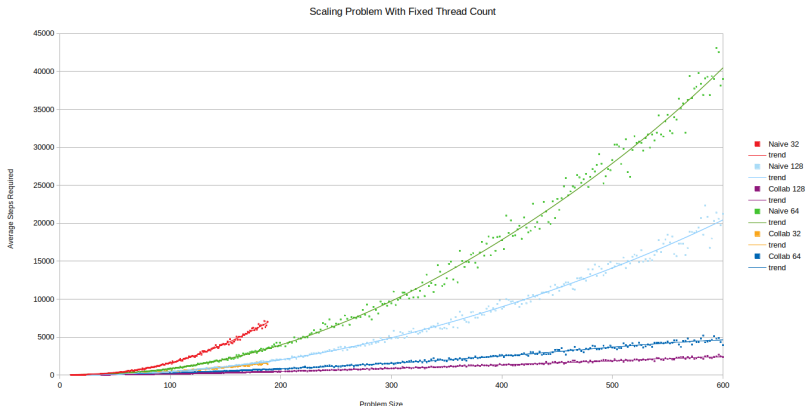


## Experimental: Counting Steps

- ▶ We utilize a global step counter, that counts the average number of increment operations across all threads
- ▶ This allows us to compare the scaling behavior of naive RIC compared to collaborative RIC in regard to the problem size, regardless of the specific hardware used
- ▶ Because both algorithms are random in nature, each data-point is the mean value of 250-500 random samples
- ▶ We test various problem sizes to compare the scaling behavior

# Scaling the problem size

- ▶ The upper curves are naive execution for 32, 64 and 128 threads for problems up to 600 queens
- ▶ The lower curves are collaborative execution



## Experiment interpretation

- ▶ All collaborative curves show much better scaling behavior
- ▶ Using more threads always reduces the required steps for both versions
- ▶ 32-threaded collaborative execution is still preferable to even 128-threaded naive execution
- ▶ The collaborative curves appear almost linear. Likely to be the shallow beginnings of a higher-order curve
- ▶ Collaborative execution is clearly superior
- ▶ Further experiments concerning high thread counts, real-time and alternative implementations are in the thesis

## Further scaling

- ▶ The prior implementation is limited to execution on a CPU
- ▶ We observed a significant speedup for both versions by using a high thread count, even far beyond the CPU hardware capabilities
- ▶ It could be nice if we could scale the thread count even higher
- ▶ For this, we use a GPU, which has far more computing units
- ▶ However, each unit is slower than a CPU core, and has other special restrictions that may destroy any speedup gained from increased thread count

## GPU considerations: Warps

- ▶ GPU threads operate best when they compute data independently from other threads
- ▶ Warps are groups of usually 32 threads, that must always execute the same instruction at the same time, though they may operate on different data
- ▶ Branches in code force execution within a warp to be sequential
- ▶ Both RIC versions operate mostly independently from other threads, but there are many branches within the central loop and the increment operation


## GPU considerations: Memory

- ▶ Solutions are too big to be stored locally in registers of a thread or in shared memory
- ▶ ... so we store them in global memory
- ▶ Global memory access is comparatively slow
- ▶ SIMD operations allow for access to 32 Bytes of continuous memory that is distributed across a warp
- ▶ This requires threads within a warp to access continuous memory in the same instruction
- ▶ If the solutions of each thread are stored in continuous memory, a warp would access 32 disparate memory locations
- ▶  $\Rightarrow$  The solutions for all threads must be stored in an interleaved manner


# Coalescing global memory access

4 Threads in a warp attempt to access 3 queens

Each Solution is continuous. No memory access coalescing



Solutions are interleaved. Warps coalesce memory access



## GPU considerations: restricting freedom

- ▶ The lock-based freedom-restrictions can not be used efficiently on a GPU
- ▶ The coupled nature of a warp would force all threads within a warp to block if any of them is copied from
- ▶ CUDA also offers no build-in lock implementation
- ▶  $\Rightarrow$  We must use a barrier-based implementation



## Barriers: Phases


- ▶ All threads must reach a barrier before any of them continue execution
- ▶ Split execution into three distinct phases that are separated by barriers
- ▶ Phase 1: All threads increment their solutions
- ▶ Phase 2: Compute, whether all of the threads failed
- ▶ Phase 3:
  - ▶ If all threads failed: Reset all threads
  - ▶ If not all threads failed: copy from active to inactive threads
- ▶ Repeat until a solution is found

## Barrier RIC pseudo-code

```
1 execute in parallel with shared variable failAll
2 {
3     solution = empty solution;
4     while(failAll)    //Repeat until success
5     {
6         for i=row to N    //increment N times
7         {
8             barrier;
9             extended = increment(solution , i);
10            failAll = true
11            barrier;
12            if(extended)    //This thread extended
13                failAll = false;
14            barrier;
15            if(failAll) //All threads failed
16                {
17                    reset(solution); //Reset all threads
18                    break; //Repeat from empty solutions
19                }
20            else if(!extended) //Not all threads failed
21                copy(solution); //Copy from other threads
22        }}
```

# A GPU test

- ▶ Slower execution times force us to test fewer values



## GPU test interpretation

- ▶ All collaborative data-series still perform better than their naive counterparts on a GPU
- ▶ Comparatively low-threaded naive execution (256 threads) performs extremely badly
- ▶ ... while collaborative data-series are closer together
- ▶ A 64-threaded CPU barriers implementation performs slightly better than all GPU versions
- ▶ The fact that the results are comparable is great
- ▶ GPU execution is unlikely to be better than CPU execution but might be, given the right hardware

# Conclusion

- ▶ RIC is an easy-to-implement, general algorithm for solving CSPs
- ▶ RIC can be sped up by massively parallel execution
- ▶ Collaborative RIC is an improved, just as general, and almost as easy to implement version of RIC
- ▶ Collaborative RIC is preferable to naive RIC on a CPU as well as a GPU
- ▶ CPU execution is likely preferable to GPU execution, but might not be, dependent on the available hardware

Questions?

Questions?